

# CS6322: Information Retrieval

## Sanda Harabagiu

### Lecture 8: Web search basics



# Brief (non-technical) history

---

- Early keyword-based engines ca. 1995-1997
  - Altavista, Excite, Infoseek, Inktomi, Lycos
- Paid search ranking: Goto (morphed into Overture.com → Yahoo!)
  - Your search ranking depended on how much you paid
  - Auction for keywords: **casino** was expensive!

# Brief (non-technical) history

---

- 1998+: Link-based ranking pioneered by Google
  - Blew away all early engines save Inktomi
  - Great user experience in search of a business model
  - Meanwhile Goto/Overture's annual revenues were nearing \$1 billion
- Result: Google added paid search “ads” to the side, independent of search results
  - Yahoo followed suit, acquiring Overture (for paid placement) and Inktomi (for search)
- 2005+: Google gains search share, dominating in Europe and very strong in North America
  - 2009: Yahoo! and Microsoft propose combined paid search offering

nigritude ultramarine - Google Search - Mozilla Firefox

File Edit View Go Bookmarks Yahoo! Tools Help

http://www.google.com/search?hl=en&q=nigritude+ultramarine&btnG=Google+Search

Getting Started Latest Headlines

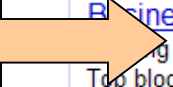
pragh60@gmail.com | My Account | Sign out

Google Web Images Groups News Froogle Local more »


nigritude ultramarine Search Advanced Search Preferences

Web Results 1 - 10 of about 185,000 for **nigritude ultramarine**. (0.35 seconds)

**Paid Search Ads**



**Algorithmic results.**



**Anil Dash: Nigritude Ultramarine**  
Do me a favor: Link to this post with the phrase **Nigritude Ultramarine**. ... Just placed a link to your **Nigritude Ultramarine** article on my weblog. Cheers! ...  
[www.dashes.com/anil/2004/06/04/nigritude\\_ultra](http://www.dashes.com/anil/2004/06/04/nigritude_ultra) - 101k - Mar 1, 2006 -  
[Cached](#) - [Similar pages](#)

**Nigritude Ultramarine FAQ**  
**Nigritude Ultramarine** FAQ - frequently asked questions about **nigritude ultramarine** and the realted SEO contest.  
[www.nigritudeultramarines.com/](http://www.nigritudeultramarines.com/) - 59k - [Cached](#) - [Similar pages](#)

**SEO contest - Wikipedia, the free encyclopedia**  
The **nigritude ultramarine** competition by SearchGuild is widely acclaimed as ...  
Comparison of search results for **nigritude ultramarine** during and after the ...  
[en.wikipedia.org/wiki/Nigritude\\_ultramarine](http://en.wikipedia.org/wiki/Nigritude_ultramarine) - 37k - [Cached](#) - [Similar pages](#)

**Slashdot | How To Get Googled, By Hook Or By Crook**  
The current 3rd result showcases the "**Nigritude Ultramarine** Fighting Force" who ... When discussing **nigritude ultramarine** [slashdot.org] it is important to ...  
[slashdot.org/article.pl?sid=04/05/09/1840217](http://slashdot.org/article.pl?sid=04/05/09/1840217) - 110k - [Cached](#) - [Similar pages](#)

**The Nigritude Ultramarine Search Engine Optimization Contest**  
It's sweeping the web -- or at least search engine optimizers -- a new contest to rank tops for the term **nigritude ultramarine** on Google.  
[searchenginewatch.com/sereport/article.php/3360231](http://searchenginewatch.com/sereport/article.php/3360231) - 57k - [Cached](#) - [Similar pages](#)

**Sponsored Links**

**Business Blogging Seminar**  
g to L.A. March 16  
Top bloggers reveal key techniques  
[www.blogbusinesssummit.com](http://www.blogbusinesssummit.com)  
Los Angeles, CA

**Full-Time SEO & SEM Jobs**  
Find companies big & small hiring full-time SEO & SEM pros right now  
[CareerBuilder.com](http://CareerBuilder.com)


**SEO Contests**  
Information on SEO Contests like the **Nigritude Ultramarine** contest.  
[www.seo-contests.com/](http://www.seo-contests.com/)

**The SEO Book**  
**Nigritude Ultramarine** & SEO secrets  
Fun, free, raw, & different.  
[www.seobook.com](http://www.seobook.com)

Music - Dance - Electronic  
Overstock.com

Done

# Web search basics



# User Needs

---

- Need [Brod02, RL04]
  - Informational – want to learn about something (~40% / 65%)

Low hemoglobin
  - Navigational – want to go to that page (~25% / 15%)

United Airlines
  - Transactional – want to do something (web-mediated) (~35% / 20%)
    - Access a service

Seattle weather
    - Downloads


Mars surface images
    - Shop

Canon S410
  - Gray areas
    - Find a good hub

Car rental Brasil
    - Exploratory search “see what’s there”

# How far do people look for results?

“When you perform a search on a search engine and don't find what you are looking for, at what point do you typically either revise your search, or move on to another search engine? (Select one)”



(Source: [iprospect.com](http://iprospect.com) WhitePaper\_2006\_SearchEngineUserBehavior.pdf)

# Users' empirical evaluation of results

---

- Quality of pages varies widely
  - Relevance is not enough
  - Other desirable qualities (non IR!!)
    - Content: Trustworthy, diverse, non-duplicated, well maintained
    - Web readability: display correctly & fast
    - No annoyances: pop-ups, etc
- Precision vs. recall
  - On the web, recall seldom matters
- What matters
  - Precision at 1? Precision above the fold?
  - Comprehensiveness – must be able to deal with obscure queries
    - Recall matters when the number of matches is very small
- User perceptions may be unscientific, but are significant over a large aggregate




# Users' empirical evaluation of engines

---

- Relevance and validity of results
- UI – Simple, no clutter, error tolerant
- Trust – Results are objective
- Coverage of topics for polysemic queries
- Pre/Post process tools provided
  - Mitigate user errors (auto spell check, search assist,...)
  - Explicit: Search within results, more like this, refine ...
  - Anticipative: related searches
- Deal with idiosyncrasies
  - Web specific vocabulary
    - Impact on stemming, spell-check, etc
  - Web addresses typed in the search box
- “The first, the last, the best and the worst ...”

# The Web document collection



- No design/co-ordination
- Distributed content creation, linking, democratization of publishing
- Content includes truth, lies, obsolete information, contradictions ...
- Unstructured (text, html, ...), semi-structured (XML, annotated photos), structured (Databases)...
- Scale much larger than previous text collections ... but corporate records are catching up
- Growth – slowed down from initial “volume doubling every few months” but still expanding
- Content can be *dynamically generated*

# Spam

---

- (Search Engine Optimization)

# The trouble with paid search ads ...

---

- It costs money. What's the alternative?
- *Search Engine Optimization:*
  - “Tuning” your web page to rank highly in the algorithmic search results for select keywords
  - Alternative to paying for placement
  - Thus, intrinsically a marketing function
- Performed by companies, webmasters and consultants (“Search engine optimizers”) for their clients
- Some perfectly legitimate, some very shady

# Search engine optimization (Spam)


---

- Motives
  - Commercial, political, religious, lobbies
  - Promotion funded by advertising budget
- Operators
  - Contractors (Search Engine Optimizers) for lobbies, companies
  - Web masters
  - Hosting services
- Forums
  - E.g., Web master world ( [www.webmasterworld.com](http://www.webmasterworld.com) )
    - Search engine specific tricks
    - Discussions about academic papers 😊

# Simplest forms

- First generation engines relied heavily on *tf/idf*
  - The top-ranked pages for the query **maui resort** were the ones containing the most **maui**'s and **resort**'s
- SEOs responded with dense repetitions of chosen terms
  - e.g., **maui resort maui resort maui resort**
  - Often, the repetitions would be in the same color as the background of the web page
    - Repeated terms got indexed by crawlers
    - But not visible to humans on browsers

Pure word density cannot  
be trusted as an IR signal



# Variants of keyword stuffing

---

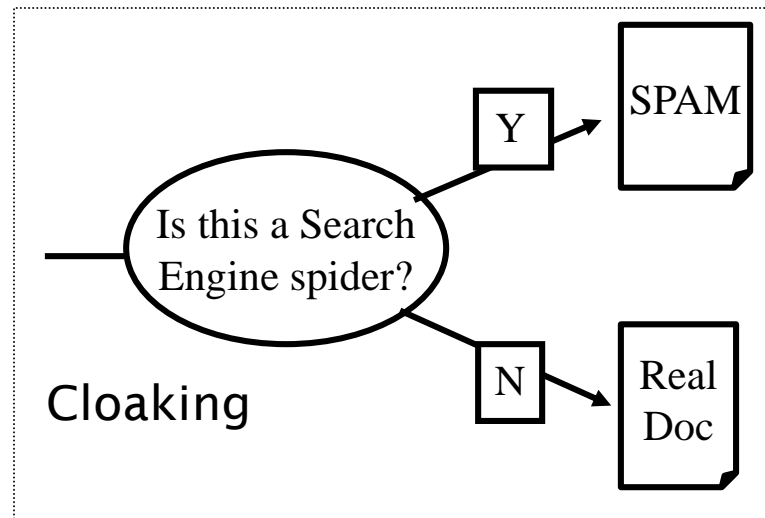
- Misleading meta-tags, excessive repetition
- Hidden text with colors, style sheet tricks, etc.

## **Meta-Tags =**

"... London hotels, hotel, holiday inn, hilton, discount, booking, reservation, sex, mp3, britney spears, viagra, ..."

# Cloaking

- Serve fake content to search engine spider
- DNS cloaking: Switch IP address. Impersonate





# More spam techniques

---

- **Doorway pages**

- Pages optimized for a single keyword that re-direct to the real target page

- **Link spamming**

- Mutual admiration societies, hidden links, awards – more on these later
- *Domain flooding*: numerous domains that point or re-direct to a target page

- **Robots**

- Fake query stream – rank checking programs
  - “Curve-fit” ranking programs of search engines
- Millions of submissions via Add-Url

# The war against spam

---

- Quality signals - Prefer authoritative pages based on:
  - Votes from authors (linkage signals)
  - Votes from users (usage signals)
- Policing of URL submissions
  - Anti robot test
- Limits on meta-keywords
- Robust link analysis
  - Ignore statistically implausible linkage (or text)
  - Use link analysis to detect spammers (guilt by association)
- Spam recognition by machine learning
  - Training set based on known spam
- Family friendly filters
  - Linguistic analysis, general classification techniques, etc.
  - For images: flesh tone detectors, source text analysis, etc.
- Editorial intervention
  - Blacklists
  - Top queries audited
  - Complaints addressed
  - Suspect pattern detection

# More on spam

---

- Web search engines have policies on SEO practices they tolerate/block
  - <http://help.yahoo.com/help/us/ysearch/index.html>
  - <http://www.google.com/intl/en/webmasters/>
- Adversarial IR: the unending (technical) battle between SEO's and web search engines
- Research <http://airweb.cse.lehigh.edu/>

# Size of the web


---

# What is the size of the web ?

---

- Issues
  - The web is really infinite
    - Dynamic content, e.g., calendar
    - Soft 404: [www.yahoo.com/<anything>](http://www.yahoo.com/<anything>) is a valid page
  - Static web contains syntactic duplication, mostly due to mirroring (~30%)
  - Some servers are seldom connected
- Who cares?
  - Media, and consequently the user
  - Engine design
  - Engine crawl policy. Impact on recall.

# Nielsen NetRatings Search Engine Ratings - July 2005



The [Nielsen NetRatings MegaView Search reporting service](#) measures the search behavior of more than a million people worldwide. These web surfers have real-time meters on their computers which monitor the sites they visit.

# More statistics on searchenginewatch.com

Search Engine Watch: Tips About Internet Search Engines & Search Engine Submission - Mozilla Firefox

File Edit View Go Bookmarks Tools Help

http://searchenginewatch.com/

Getting Started Latest Headlines

The Premier Event for Search Engine Marketing & Optimization **Search Engine STRATEGIES 2005** December 5-8, 2005 • Chicago, IL

ClickZ. You are in the: ClickZ Network > ClickZ Network Navigation

**SearchEngineWatch**  
The source for search engine marketing

**Members Area With Exclusive Content**  
Already a member? > Enter Here Learn about the benefits: > Learn More

**Departments & Info**

- Home
- Latest Stories From SEW
- SEW Blog
- News From SEW & Beyond
- SEW Forums
- Come Discuss Search!
- Search Engine Submission Tips
- Web Searching Tips
- Search Engine Listings
- Search Ratings & Stats
- Search Engine Resources
- SearchDay
- Our Daily Newsletter
- Search Engine Report
- Our Monthly Newsletter
- Daily SearchCast
- Our Search Podcast
- All Newsletters & Feeds
- [XML](#) [RSS](#)
- SEW Members Area
- Exclusive Content
- Advertising Info
- More About The Site
- >> Search

**Highlights from the SEW Blog: Oct 26, 2005**  
**October 26, 2005** - Featured posts from the Search Engine Watch blog, as well as our customary search headlines from around the web.

**Featured Discussions In Our Forums**

- Google places "Dummy Bidders" in Regional Search artificially raise required bids
- A Trip Down Memory Lane...Your First Search Engine Query
- MSN AdCenter Rough Start
- Google Base: Google's Protect To Take Listings Of Anything
- More From Our Forums ...

**Shopping Search and Merchant Reputations**  
**October 25, 2005** - Consumers often look to merchant reputations on shopping search engines when deciding to make a purchase. But who determines these reputations? And can you trust them?

**Latest From the Search Engine Watch Blog**

- Google Makes Digitized Archive of Interviews with TV Legends Available Online
- Daily SearchCast On Hiatus Until Nov. 8: Webmaster Radio Hit Hard By Hurricane Wilma
- Microsoft Announces MSN Book Search: Joins Open Content Alliance
- Google's New "Travel" OneBox
- More From Our Blog ...

**Inside the Internet Movie Database**  
**October 24, 2005** - The Internet Movie Database is a standard and respected reference database for film and TV information, but for film, TV and pop-culture fans it's also a great deal of fun.

**Search Engine Forums Spotlight**  
**October 21, 2005** - Links to the week's topics from search engine forums across the web: October 2005 Google Update -

**Search Engine STRATEGIES 2005 CONFERENCE & EXPO.**  
Search Engine Watch's conference on search engine marketing comes to:

- November 29-30, 2005 Cnit Paris La Defense, Paris, France
- December 5 - 8, 2005 Chicago, IL
- February 27 - March 2, 2006 New York, NY
- March 17 - 18, 2006 Nanjing, China
- March 30 - 31, 2006 Munich, Germany
- April 25 - 26, 2006 Toronto, Canada
- April 26 - 27, 2006 Milan, Italy
- April 27-28, 2006 Tokyo, Japan
- May 31 - June 2, 2006 London England
- July 10-11, 2006 Miami, Florida
- August 7 - 10, 2006 San Jose, CA

**>> Search Engine Watch Marketplace**


- WebPosition™**  
Optimize your site and drive qualified visitors.
- BRUCECLAY.COM**  
866-517-1900  
Search Engine Optimization Tools and Services
- InfoSearch Media**  
Professional, Custom Content
- Trellian**  
Keyword Research & Competitive Intelligence
- morevisibility**  
Free Search Marketing Report  
Get Your Site More Visibility
- YAHOO! SEARCH MARKETING**  
Reach 80% of the web on Yahoo! and more
- Text Link Ads**  
Increase Traffic & Search Engine Rankings
- WebSourced's KeywordRanking.com**  
Search Engine Marketing Services Free Report

Done

# Web Challenges

---

- Distributed data
- High percentage of volatile data
- Large volume
- Unstructured and redundant data
- Quality of data
- Heterogeneous data





Several architectures  
for search engines

➔ difference from standard IR systems: all queries must be answered without accessing the text.



# Centralized Architecture





**Gatherers** = collect and extract indexing information from one or more web servers

**Brokers** = provide the indexing information and the query interface to the data gathered

# Search Engines

---

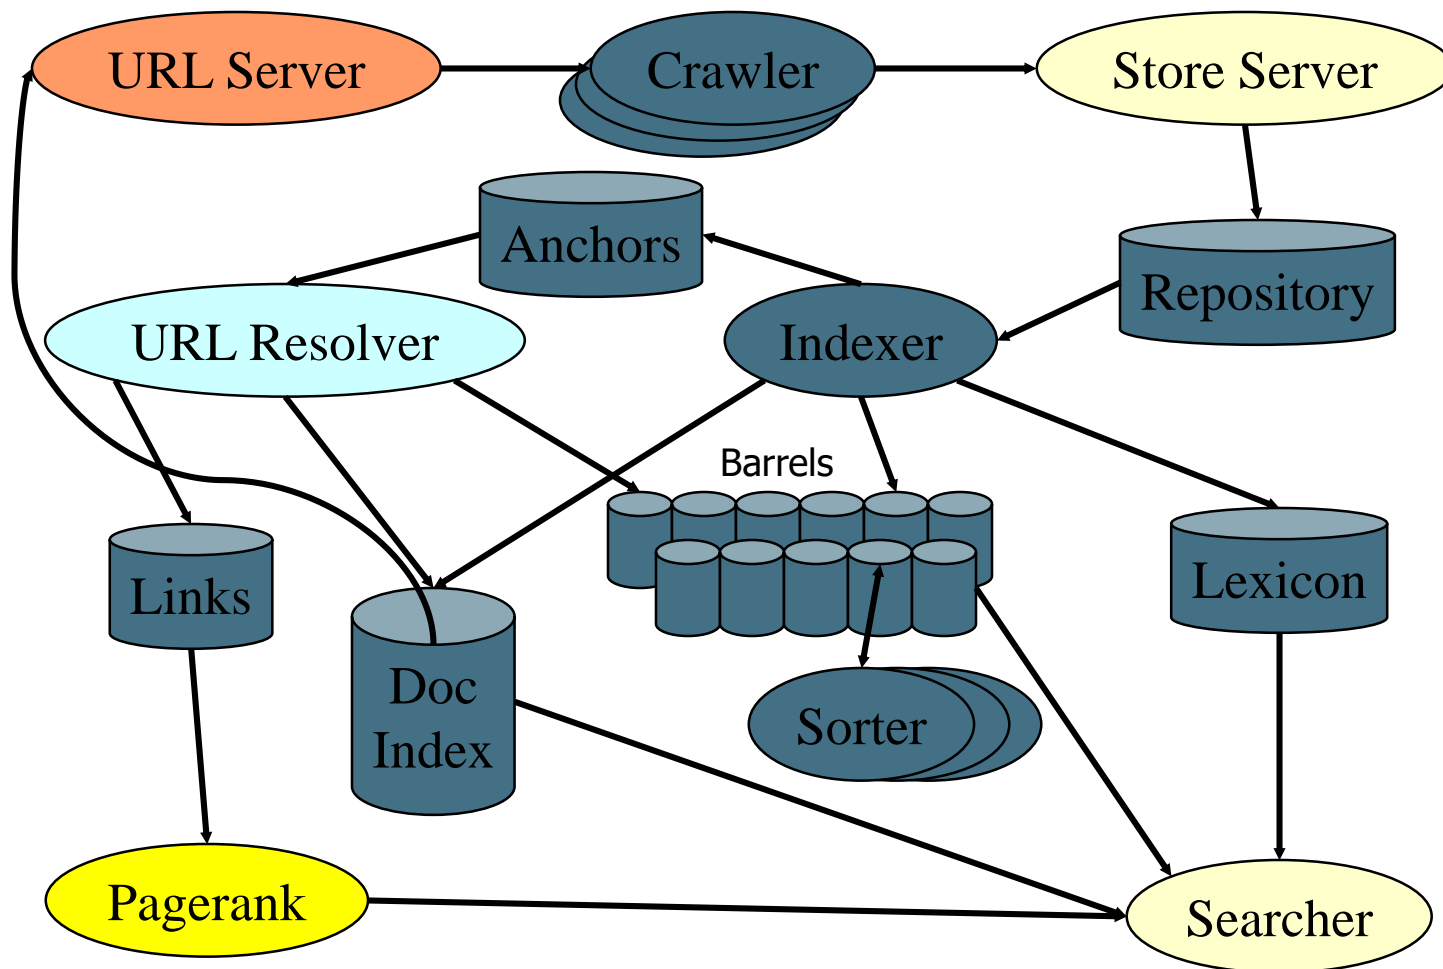
- Google case (1997)
  - Name is a common spelling of *googol*  $10^{100}$
  - Goal: build a very large-scale search engine
  - Index : 1997 100million; 2002 2 billion
  - In 1997 their forecast was that in 2000 they will have 1 billion page index (they had 2 billion pages!!!)
- World Wide Web Worm (1994)
  - Index 1994 110,000 web pages

# Web Search Engines: Google

---

- *Hyperlinks* represent new forms of information
- *Hypertexts* have existed and been studied for a long time
- New → a large number of hyperlinks created by independent individuals
- Hyperlinks provide a valuable source of information for Web IR → Link analysis

# Google Anatomy




# Google Indexer

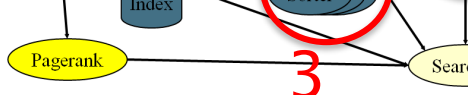
## Functions

- Reads the repository
- Uncompress the documents, parse them
- Each document is converted into a set set of word occurrences → hits
  - 1 hit records [word, position in document, font size, capitalization]
- The hits are distributed in a set of “barrels” partially sorted forward index
- Parses all links and stores info about them in anchors file:
  - <where the links point from, to, text of link>

## Google Anatomy




# Google Anatomy

- Converts relative URLs in absolute URLs
    - Translate them in docIDs
  - Puts the anchor text into the forward index, associated with the docID that the anchor points to
  - Generates a database of links = pairs of docIDs
- 
- A diagram illustrating a search engine architecture. It features several components: a blue cylinder labeled 'Index', a yellow oval labeled 'Pagerank', a yellow oval labeled 'Search', and a red circle labeled '3'. The 'Index' component is connected to the 'Pagerank' component. The 'Pagerank' component is connected to the 'Search' component. The 'Search' component is also connected to the 'Index' component. The red circle '3' is positioned between the 'Index' and 'Search' components, indicating a specific step or process in the search engine's operation.

**Sorter** – takes the barrels (sorted by **docID**) and resorts them by **wordID** to generate the inverted index.

- Sorting in place → little temporary space is needed.

**Dump Lexicon** – a program that takes the inverted list and the lexicon produced by the indexer → generates new lexicon from searcher.



# Major Data Structures in Google


Optimized for crawling/indexing/searching large document collections.

Trick → avoid disk seeks (10ms per seek)

→ influences the design of data structures

- Big files
- Repository
- Document Index
- Lexicon
- Hit list
- Forward index
- Inverted index

## Google Anatomy





# Big Files

---

- Virtual files spanning multiple file systems → addressable by 64 bit integers
- Allocation among multiple file systems is handled automatically
  - Allocates and de-allocates descriptors
  - Compression options


# Repository

- Contains the full HTML of every web page
- Compression using zlib

Repository: 53.5 GB = 147.8 GB uncompressed

sync	length	Compressed packet
------	--------	-------------------


sync	length	Compressed packet
------	--------	-------------------



...

docid	ecode	urlen	pagelen	url	page
-------	-------	-------	---------	-----	------

Packet (stored compressed in repository)




- Each query:

current document status	Pointer to repository	Document checksum	statistics
-------------------------	-----------------------	-------------------	------------

# Lexicon: 14 million of words

- ❑ List of words + hash table of pointers
- ❑ Sanda Harabagiu research


# Google Anatomy



# Hit Lists

- Account for most of the space used in both the forward and inverted indices
- It contains 2 forms of information:
  - documents where a word occurred,
  - position in document and font info
- Encoding position, font, capitalization
  - Simple encoding (3 integers)
  - Compact encoding (hand-optimized allocation of bits)
  - Huffman encoding
- Select – hand optimized compact encoding
  - Less space than simple encoding
  - Less bit manipulation than Huffman coding

Google Anatomy



# Google Compact Coding

---

- Uses 2 bytes for every hit
- There are two types of hits: **fancy hits** and **plain hits**:
  1. Fancy lists → hits occurring in a URL, title, anchor text or meta tag
  2. Plain hits → all the other

The codification:

A plain hit has

- 1 capitalization bit
- font size (3 bits)
- 12 bits of word position in document (positions higher than 4096 are labeled 4096)
- For font size → only 7 values, since 111 signals a fancy bit

# Google Compact Coding - cont

A fancy hit has

- 1 capitalization bit
- Font size set to 7
- 4 bits encode the type of fancy hit
- 8 bits for position
- For anchor hits the 8 bits for position are split:
  - 4 bits for position anchor
  - 4 bits for hash of the **docID** the anchor occurs in

Hit: 2 bytes

Cap:1	Imp:3	Position:12		
Cap:1	Imp=7	Type:4	Position:8	
Cap:1	Imp=7	Type:4	Hash:4	Pos:4

# Forward Index

- It is partially sorted
- Stored in 64 barrels. Each **barrel** holds a range of word Ids
- If a document contains words that fall in a particular barrel, the **docID** is recorded into the barrel followed by a list of word Ids with hit-lists corresponding to the words


Forward barrels: total 43 GB

docid	Wordid:24	Nhits:8	Hit hit hit hit
	Wordid:24	Nhits:8	Hit hit hit hit
	Null wordid		
docid	Wordid:24	Nhits:8	Hit hit hit hit
	Wordid:24	Nhits:8	Hit hit hit hit
	Wordid:24	Nhits:8	Hit hit hit hit
	Null wordid		

# Forward Index -cont

- **DocIds** are duplicated –reasonable for 64 buckets but saves considerable time and coding complexity
- Instead of storing actual word Id, store relative difference to minimum word ID in the barrel. Use only 24 bits for word ID.
- 8 bits → hit list length

Google Anatomy





# Inverted Index

---

- On the same barrels as forward index
- For every valid **wordID** – there is a pointer in the lexicon to the barrel that the **wordID** falls into:
  - Barrel points to a doclist of **docIDs** + corresponding hit lists.
- Issue: in what order the **docIDs** should appear in the doclist ?
  - **Solution 1**: store them sorted by docID – quick mapping for multiple word queries
  - **Solution 2**: store them sorted by a ranking of occurrence of the word in each document – merging is difficult,
  - **Solution (Google)**: compromise: 2 sets of inverted barrels:
    - 1 set for hit lists which include title and anchor hits
    - 1 set for the other hit lists.
    - Check the first set – if not enough matches, check the 2<sup>nd</sup> set

# Crawling the Web

---

## Running a Web Crawler is a challenging task

- Crawling is the most fragile application – it involves interacting with hundreds of thousands of Web servers and various name servers → beyond the control of the system
- Google solution: fast distributed crawling system
  - The URL server serves lists of URLs to 3-5 crawlers (both server and crawlers are implemented in Python)
  - Each crawler keeps 300 connections open at once to retrieve Web pages at a fast enough pace.
  - 100 web pages/sec  $\approx$  600k/sec

# Google crawler

- Performance stress: DNS lookup
- Each crawler maintains its own DNS cache (it does not need to do a DNS lookup before crawling each document)
- Each of the 300 connections can be in a different state: Lookup DNS, Connect to host, Send request or Receive response

- **The Robot Exclusion Standard**

(devised in 1994 by Martin Kostner)

- Declares that a Web server administrator should create a document accessible at the relative URL /robots.txt

- E.g.


User-agent: \*

Disallow: /cgi

Disallow: /stats

Try: [www.utdallas.edu/robots.txt](http://www.utdallas.edu/robots.txt)

Google Anatomy



# The robots.txt file

---

- 3 basic directives can be in robots.txt: **User-agent**, **Allow** and **Disallow**
- If the robots.txt file contains:

User-agent: \*

Disallow: /

the administrator wants to shut out all robots from the entire site

- Multiple User-agents can be specified:

User-agent: friendly indexes

User-agent: search-thingy

Disallow: /cgi-bin/

Allow: /

- Another example:

User-agent: \*

Disallow: /

User-agent: search-thingy

Allow: /

➔ all robots should go away, except the search-thingy robot


# Google Initial Performance

Storage Statistics	
Total size of Fetched Pages	147.8 GB
Compressed Repository	53.5 GB
Short Inverted Index	4.1 GB
Full Inverted Index	37.2 GB
Lexicon	293 MB
Temporary anchor data (not in total)	6.6 GB
Document Index Incl. Variable Width Data	9.7 GB
Links Database	3.9 GB
Total Without Repository	55.2 GB
Total With Repository	108.7 GB

Web Page Statistics	
Number of Web Pages Fetched	24 million
Number of URLs seen	76.5 million
Number of Email addresses	1.7 million
Number of 404's	1.6 million

# The Web Macroscopic Structure

- AltaVista experiments



90% are in a single, giant Strongly Connected Component (SCC)

IN = pages that can reach SCC, but cannot be reached from it

OUT = pages reached from SCC, but cannot access SCC

TENDRILS = pages that cannot reach SCC and cannot be reached by SCC

All 4 sets have roughly the same size → big surprise !

# More Data

---

- The diameter of the central core of SCC is at least 28, the diameter of the graph is  $> 500$
- The probability that a path exists from the source to the destination 0.24
  - If a directed path exists, its average length is 16
- The Power law for in-degree
  - The probability that a node has in-degree  $l$  is proportional to  $1/l^x$  for some  $x > 1$ 
    - $x = 2.1$  in the AltaVista May '99 Crawl

# What can we attempt to measure?

---

- The relative sizes of search engines
  - The notion of a page being indexed is still *reasonably* well defined.
  - Already there are problems
    - Document extension: e.g. engines index pages not yet crawled, by indexing anchor text.
    - Document restriction: All engines restrict what is indexed (first  $n$  words, only relevant words, etc.)
- The coverage of a search engine relative to another particular crawling process.



# New definition?


---

(IQ is whatever the IQ tests measure.)

- The statically indexable web is whatever search engines index.
- Different engines have different preferences
  - max url depth, max count/host, anti-spam rules, priority rules, etc.
- Different engines index different things under the same URL:
  - frames, meta-keywords, document restrictions, document extensions, ...

# Relative Size from Overlap

## Given two engines A and B



**Sample** URLs randomly from A

**Check** if contained in B and vice versa

$$A \cap B = (1/2) * \text{Size A}$$

$$A \cap B = (1/6) * \text{Size B}$$

$$(1/2) * \text{Size A} = (1/6) * \text{Size B}$$

$$\therefore \text{Size A} / \text{Size B} =$$

$$(1/6) / (1/2) = 1/3$$

**Each test involves:** (i) Sampling (ii) Checking

# Sampling URLs

---

- Ideal strategy: Generate a random URL and check for containment in each index.
- Problem: Random URLs are hard to find! Enough to generate a random URL contained in a given Engine.
- Approach 1: Generate a random URL contained in a given engine
  - Suffices for the estimation of relative size
- Approach 2: Random walks / IP addresses
  - In theory: might give us a true estimate of the size of the web (as opposed to just relative sizes of indexes)

# Statistical methods

---

- Approach 1
  - Random queries
  - Random searches
- Approach 2
  - Random IP addresses
  - Random walks

# Random URLs from random queries

- Generate random query: how?
  - **Lexicon**: 400,000+ words from a web crawl
  - **Conjunctive Queries**:  $w_1$  and  $w_2$   
*e.g., vocalists AND rsi*
- Get 100 result URLs from engine A
- Choose a random URL as the candidate to check for presence in engine B
- This distribution induces a probability weight  $W(p)$  for each page.
- Conjecture:  $W(SE_A) / W(SE_B) \sim |SE_A| / |SE_B|$

Not an English dictionary

# Query Based Checking

---

- *Strong Query* to check whether an engine  $B$  has a document  $D$ :
  - Download  $D$ . Get list of words.
  - Use 8 low frequency words as AND query to  $B$
  - Check if  $D$  is present in result set.
- Problems:
  - Near duplicates
  - Frames
  - Redirects
  - Engine time-outs
  - Is 8-word query good enough?

# Advantages & disadvantages

---

- Statistically sound under the induced weight.
- Biases induced by random query
  - Query Bias: Favors content-rich pages in the language(s) of the lexicon
  - Ranking Bias: *Solution*: Use conjunctive queries & fetch all
  - Checking Bias: Duplicates, impoverished pages omitted
  - Document or query restriction bias: engine might not deal properly with 8 words conjunctive query
  - Malicious Bias: Sabotage by engine
  - Operational Problems: Time-outs, failures, engine inconsistencies, index modification.

# Random searches

---

- Choose random searches extracted from a local log [Lawrence & Giles 97] or build “random searches” [Notess]
  - Use only queries with small result sets.
  - Count normalized URLs in result sets.
  - Use ratio statistics



# Advantages & disadvantages

---

- Advantage
  - Might be a better reflection of the human perception of coverage
- Issues
  - Samples are correlated with source of log
  - Duplicates
  - Technical statistical problems (must have non-zero results, ratio average not statistically sound)

# Random searches

---

- 575 & 1050 queries from the NEC RI employee logs
- 6 Engines in 1998, 11 in 1999
- Implementation:
  - Restricted to queries with  $< 600$  results in total
  - Counted URLs from each engine after verifying query match
  - Computed size ratio & overlap for individual queries
  - Estimated index size ratio & overlap by averaging over all queries

# Queries from Lawrence and Giles study

---

- *adaptive access control*
- *neighborhood preservation topographic*
- *hamiltonian structures*
- *right linear grammar*
- *pulse width modulation neural*
- *unbalanced prior probabilities*
- *ranked assignment method*
- *internet explorer favourites importing*
- *karvel thornber*
- *zili liu*
- *softmax activation function*
- *bose multidimensional system theory*
- *gamma mlp*
- *dvi2pdf*
- *john oliensis*
- *rieke spikes exploring neural*
- *video watermarking*
- *counterpropagation network*
- *fat shattering dimension*
- *abelson amorphous computing*

# Random IP addresses

---

- Generate random IP addresses
- Find a web server at the given address
  - If there's one
- Collect all pages from server
  - From this, choose a page at random

# Random IP addresses

---

- HTTP requests to random IP addresses
  - Ignored: empty or authorization required or excluded
  - [Lawr99] Estimated 2.8 million IP addresses running crawlable web servers (16 million total) from observing 2500 servers.
  - OCLC using IP sampling found 8.7 M hosts in 2001
    - Netcraft [Netc02] accessed 37.2 million hosts in July 2002
- [Lawr99] exhaustively crawled 2500 servers and extrapolated
  - Estimated size of the web to be 800 million pages
  - Estimated use of metadata descriptors:
    - Meta tags (keywords, description) in 34% of home pages, Dublin core metadata in 0.3%

# Advantages & disadvantages

---

- Advantages
  - Clean statistics
  - Independent of crawling strategies
- Disadvantages
  - Doesn't deal with duplication
  - Many hosts might share one IP, or not accept requests
  - No guarantee all pages are linked to root page.
    - Eg: employee pages
  - Power law for # pages/hosts generates bias towards sites with few pages.
    - But bias can be accurately quantified IF underlying distribution understood
  - Potentially influenced by spamming (multiple IP's for same server to avoid IP block)

# Random walks

---

- View the Web as a directed graph
- Build a random walk on this graph
  - Includes various “jump” rules back to visited sites
    - Does not get stuck in spider traps!
    - Can follow all links!
  - Converges to a stationary distribution
    - Must assume graph is finite and independent of the walk.
    - Conditions are not satisfied (cookie crumbs, flooding)
    - Time to convergence not really known
  - Sample from stationary distribution of walk
  - Use the “strong query” method to check coverage by SE

# Advantages & disadvantages

---

- Advantages
  - “Statistically clean” method at least in theory!
  - Could work even for infinite web (assuming convergence) under certain metrics.
- Disadvantages
  - List of seeds is a problem.
  - Practical approximation might not be valid.
  - Non-uniform distribution
    - Subject to link spamming



# Conclusions

---

- No sampling solution is perfect.
- Lots of new ideas ...
- ....but the problem is getting harder
- Quantitative studies are fascinating and a good research problem

# Duplicate detection

---

# Duplicate documents

---

- The web is full of duplicated content
- Strict duplicate detection = exact match
  - Not as common
- But many, many cases of near duplicates
  - E.g., Last modified date the only difference between two copies of a page

# Duplicate/Near-Duplicate Detection

---

- *Duplication*: Exact match can be detected with fingerprints
- *Near-Duplication*: Approximate match
  - Overview
    - Compute syntactic similarity with an edit-distance measure
    - Use similarity threshold to detect near-duplicates
      - E.g., Similarity > 80% => Documents are “near duplicates”
      - Not transitive though sometimes used transitively

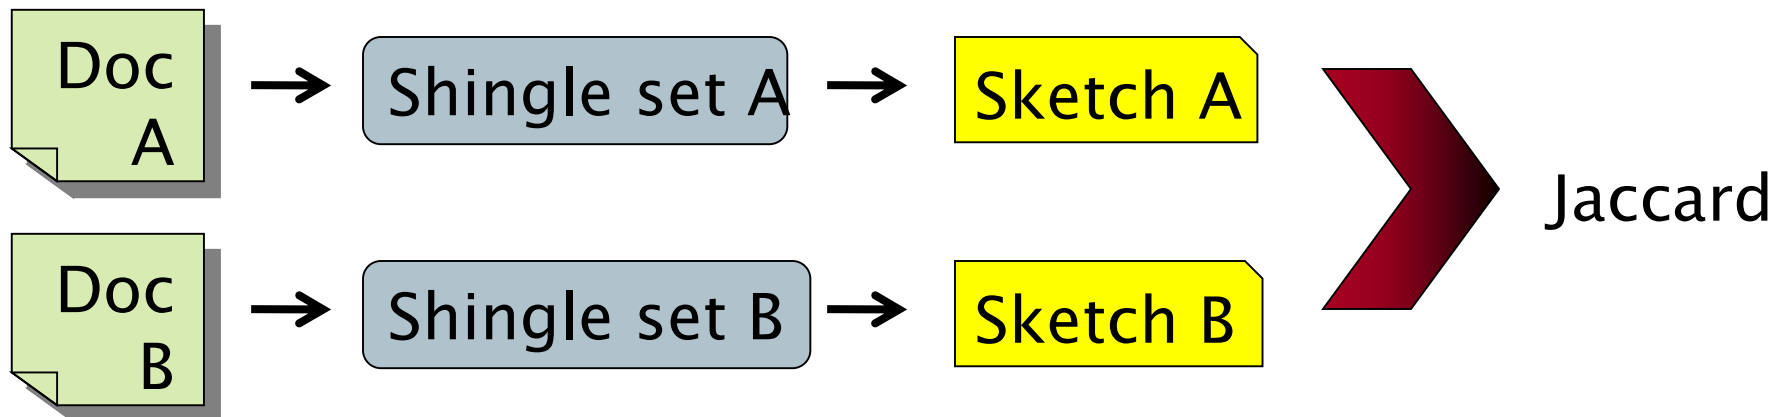
# Computing Similarity

---

- Features:
  - Segments of a document (natural or artificial breakpoints)
  - Shingles (Word N-Grams)
  - ***a rose is a rose is a rose*** →
    - a\_rose\_is\_a
    - rose\_is\_a\_rose
    - is\_a\_rose\_is
    - a\_rose\_is\_a
- Similarity Measure between two docs (= sets of shingles)
  - Set intersection
  - Specifically (Size\_of\_Intersection / Size\_of\_Union)

# Shingles + Set Intersection

- Computing exact set intersection of shingles between all pairs of documents is expensive/intractable
  - Approximate using a cleverly chosen subset of shingles from each (a *sketch*)
- Estimate  $(\text{size\_of\_intersection} / \text{size\_of\_union})$  based on a short sketch

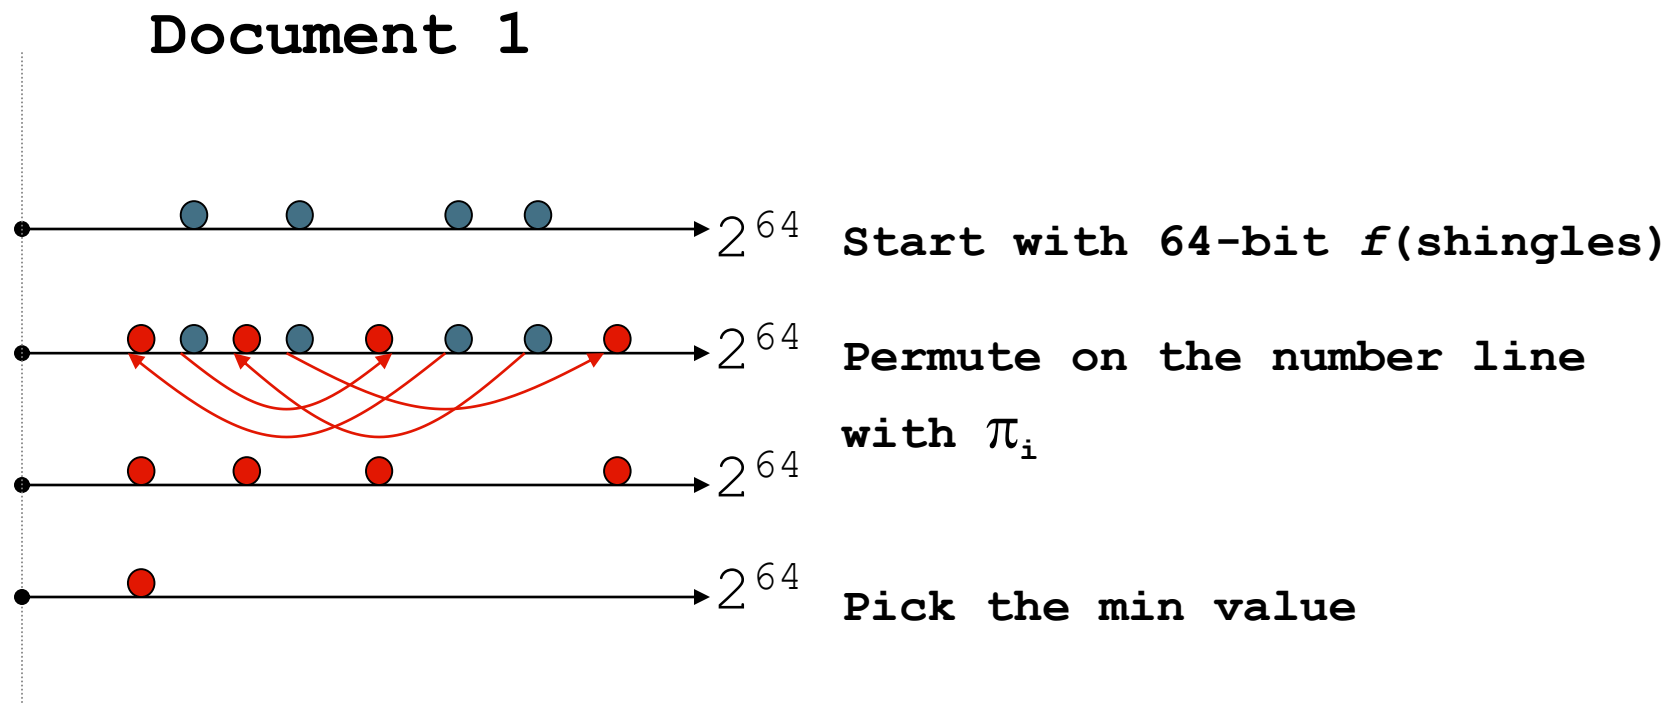


# Sketch of a document

---


- Create a “sketch vector” (of size  $\sim 200$ ) for each document
  - Documents that share  $\geq t$  (say 80%) corresponding vector elements are **near duplicates**
  - For doc  $D$ ,  $\text{sketch}_D[i]$  is as follows:
    - Let  $f$  map all shingles in the universe to  $0..2^m$  (e.g.,  $f$  = fingerprinting)
    - Let  $\pi_i$  be a *random permutation* on  $0..2^m$
    - Pick  $\text{MIN} \{ \pi_i(f(s)) \}$  over all shingles  $s$  in  $D$

# Computing Sketch[i] for Doc1






# Test if $\text{Doc1.Sketch}[i] = \text{Doc2.Sketch}[i]$



Are these equal?

Test for **200** random permutations:  $\pi_1, \pi_2, \dots, \pi_{200}$

# However...



$A = B$  iff the shingle with the MIN value in the union of Doc1 and Doc2 is common to both (i.e., lies in the intersection)

Claim: This happens with probability

$$\frac{\text{Size\_of\_intersection}}{\text{Size\_of\_union}}$$

Why?

# Set Similarity of sets $C_i, C_j$

$$\text{Jaccard}(C_i, C_j) = \frac{|C_i \cap C_j|}{|C_i \cup C_j|}$$

- View sets as columns of a matrix  $A$ ; one row for each element in the universe.  $a_{ij} = 1$  indicates presence of item  $i$  in set  $j$

- Example

$C_1$	$C_2$
0	1
1	0
1	1
0	0
1	1
0	1

$$\text{Jaccard}(C_1, C_2) = 2/5 = 0.4$$

# Key Observation

- For columns  $C_i, C_j$ , four types of rows

	$C_i$	$C_j$
A	1	1
B	1	0
C	0	1
D	0	0

- Overload notation:  $A$  = # of rows of type A
- Claim

$$\text{Jaccard}(C_i, C_j) = \frac{A}{A + B + C}$$

# “Min” Hashing

---

- Randomly **permute** rows
- **Hash**  $h(C_i)$  = index of first row with 1 in column  $C_i$
- **Surprising Property**
$$P[h(C_i) = h(C_j)] = \text{Jaccard}(C_i, C_j)$$
- **Why?**
  - Both are  $A/(A+B+C)$
  - Look down columns  $C_i, C_j$  until first **non-Type-D** row
  - $h(C_i) = h(C_j) \leftrightarrow$  type A row

# Min-Hash sketches

---

- Pick  $P$  random row permutations
- MinHash sketch

$\text{Sketch}_D =$  list of  $P$  indexes of first rows with 1 in column  $C$

- Similarity of signatures
  - Let  $\text{sim}[\text{sketch}(C_i), \text{sketch}(C_j)] =$  fraction of permutations where MinHash values agree
  - Observe  $E[\text{sim}(\text{sig}(C_i), \text{sig}(C_j))] = \text{Jaccard}(C_i, C_j)$

# Example

	$C_1$	$C_2$	$C_3$
$R_1$	1	0	1
$R_2$	0	1	1
$R_3$	1	0	0
$R_4$	1	0	1
$R_5$	0	1	0

## Signatures

	$S_1$	$S_2$	$S_3$
Perm 1 = (12345)	1	2	1
Perm 2 = (54321)	4	5	4
Perm 3 = (34512)	3	5	4

## Similarities

	1-2	1-3	2-3
Col-Col	0.00	0.60	0.4
Sig-Sig	0.00	0.67	0.00

# Implementation Trick

---

- **Permuting** universe even once is prohibitive
- **Row Hashing**
  - Pick  $P$  hash functions  $h_k: \{1, \dots, n\} \rightarrow \{1, \dots, O(n)\}$
  - **Ordering** under  $h_k$  gives random permutation of rows
- **One-pass Implementation**
  - For each  $C_i$  and  $h_k$ , keep “**slot**” for min-hash value
  - **Initialize** all  $\text{slot}(C_i, h_k)$  to **infinity**
  - **Scan rows** in arbitrary order looking for 1's
    - Suppose row  $R_j$  has 1 in column  $C_i$
    - For each  $h_k$ ,
      - if  $h_k(j) < \text{slot}(C_i, h_k)$ , then  $\text{slot}(C_i, h_k) \leftarrow h_k(j)$



# Example

	$C_1$	$C_2$
$R_1$	1	0
$R_2$	0	1
$R_3$	1	1
$R_4$	1	0
$R_5$	0	1

$$h(x) = x \bmod 5$$

$$g(x) = 2x+1 \bmod 5$$

$$h(1) = 1$$

$$g(1) = 3$$

$$h(2) = 2$$

$$g(2) = 0$$

$$h(3) = 3$$

$$g(3) = 2$$

$$h(4) = 4$$

$$g(4) = 4$$

$$h(5) = 0$$

$$g(5) = 1$$

$C_1$  slots       $C_2$  slots

1      -

3      -

1      2

3      0

1      2

2      0

1      2

2      0

1	0
---	---

2	0
---	---

# Comparing Signatures

---

- Signature Matrix  $S$ 
  - Rows = Hash Functions
  - Columns = Columns
  - Entries = Signatures
- Can compute – Pair-wise similarity of any pair of signature columns

# All signature pairs

---

- Now we have an extremely efficient method for estimating a Jaccard coefficient for a single pair of documents.
- But we still have to estimate  $N^2$  coefficients where  $N$  is the number of web pages.
  - Still slow
- One solution: locality sensitive hashing (LSH)
- Another solution: sorting (Henzinger 2006)

# More resources

---

- IIR Chapter 19